

1.

a) CQL Query:

```
DESCRIBE hoofers;
```

b) Result:

```
token@cqlsh> DESCRIBE hoofers;

CREATE KEYSPACE hoofers WITH replication = {'class': 'NetworkTopologyStrategy', 'us-east1': '3'} AND durable_writes = true;

CREATE TABLE hoofers.boats (
  bid int PRIMARY KEY,
  bname text,
  color text
) WITH additional_write_policy = '99p'
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND comment = ''
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.UnifiedCompactionStrategy'}
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND crc_check_chance = 1.0
  AND default_time_to_live = 0
  AND gc_grace_seconds = 864000
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair = 'BLOCKING'
  AND speculative_retry = '99p';

token@cqlsh> █
```

c) Answers:

The Hoofers keyspace maintains 3 copies of the data. It resides in the 'us-east1' cloud region. The Read and Write quorum sizes should each be 2 in this case. This means that at least 2 copies (a majority of 3) must accept the read or write operation to successfully occur.

2.

a) CQL CREATE Statements:

```
token@cqlsh> USE hoofers;
token@cqlsh:hoofers> CREATE TABLE Users(user_id text PRIMARY KEY, email text, joined_date date, nickname text, street text, city text, state text, zip text, genres text);
token@cqlsh:hoofers> CREATE TABLE Records(record_id text PRIMARY KEY, artist_user_id text, title text, genre text, release_date date);
token@cqlsh:hoofers> CREATE TABLE Reviews(review_id text PRIMARY KEY, user_id text, record_id text, rating int, body text, posted_at timestamp);
token@cqlsh:hoofers> CREATE TABLE Sessions(session_id text PRIMARY KEY, user_id text, record_id text, track_number int, initiate_at timestamp, leave_at timestamp, music_quality text, device text, remaining_time int, replay_count int);
token@cqlsh:hoofers> █
```

3.

a) PostgreSQL COPY commands:

```
-- COPYING to CSV file
-- users, records, reviews, and sessions
COPY Users TO '/Applications/PostgreSQL 17/Big Data HW 1/users_224p.csv' WITH (FORMAT CSV, HEADER);
COPY Records TO '/Applications/PostgreSQL 17/Big Data HW 1/records_224p.csv' WITH (FORMAT CSV, HEADER);
COPY Reviews TO '/Applications/PostgreSQL 17/Big Data HW 1/reviews_224p.csv' WITH (FORMAT CSV, HEADER);
COPY Sessions TO '/Applications/PostgreSQL 17/Big Data HW 1/sessions_224p.csv' WITH (FORMAT CSV, HEADER);
```

4.

a) First CQL Query:

```
select record_id, genre, title
from records
where artist_user_id = 'user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3';
```

b) Result:

```
token@cqlsh:hoofers> select record_id, genre, title from records where artist_user_id = 'user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3';
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
token@cqlsh:hoofers> █
```

c) Modified CQL Query:

```
SELECT record_id, genre, title
FROM records
WHERE artist_user_id = 'user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3'
ALLOW FILTERING;
```

b) Result:

```
token@cqlsh:hoofers> select record_id, genre, title from records where artist_user_id = 'user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3' ALLOW FILTERING;
```

record_id	genre	title
record_62389a63-e95f-43d1-acea-aalbac0e0050	Gospel	Result guess for
record_116fbdd6-e706-41f7-9809-12e174e48e8f	Jazz	Discover rate
record_2406e933-23e3-4db1-acf9-3c863d48bff6	Country	General job heavy
record_3e4ed054-cf1a-4a04-8e97-e0177c6d3575	Folk	Summer civil political beat
record_91c6325d-b17f-4f4c-be6d-3517b2173a9f	Country	Statement matter
record_57061d35-de20-4bf1-9aac-a689f0db7e16	Soul	Would determine
record_822961a3-946a-49ff-8173-74d4035286b9	Gospel	Apply size
record_cbf93efd-2deb-48ae-ad73-83aa088c6f13	Soul	Democratic what
record_5cbf14c7-7b54-4e32-bfce-cba507c7277f	Jazz	Bar talk long

(9 rows)
token@cqlsh:hoofers> █

5.

a) CQL Create Statement:

```
CREATE TABLE records_q5(artist_user_id text, record_id text, genre text, release_date date, title text,  
PRIMARY KEY (artist_user_id, record_id));
```

b) CQL Query:

```
SELECT record_id, genre, title
```

```
FROM records_q5
```

```
WHERE artist_user_id = 'user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3';
```

c) Result:

```
token@cqlsh:hoofers> select record_id, genre, title from records_q5 where artist_user_id = 'user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3';
```

record_id	genre	title
record_116fbdd6-e706-41f7-9809-12e174e48e8f	Jazz	Discover rate
record_2406e933-23e3-4db1-acf9-3c863d48bff6	Country	General job heavy
record_3e4ed054-cf1a-4a04-8e97-e0177c6d3575	Folk	Summer civil political beat
record_57061d35-de20-4bf1-9aac-a689f0db7e16	Soul	Would determine
record_5cbf14c7-7b54-4e32-bfce-cba507c7277f	Jazz	Bar talk long
record_62389a63-e95f-43d1-acea-aalbac0e0050	Gospel	Result guess for
record_822961a3-946a-49ff-8173-74d4035286b9	Gospel	Apply size
record_91c6325d-b17f-4f4c-be6d-3517b2173a9f	Country	Statement matter
record_cbf93efd-2deb-48ae-ad73-83aa088c6f13	Soul	Democratic what

(9 rows)
token@cqlsh:hoofers>

d) Explanation:

Using the partition key created partitions on the artist_user_id ran without any error from Cassandra. This is because Cassandra no longer has to scan all the partitions to find the relevant data, so we have greater efficiency in retrieving the records. Therefore, there is no error or warning from Cassandra. We include record_id in the primary key so that we can uniquely identify each record given the artist_user_id and record_id.

6.

a) CQL Query:

```
SELECT record_id, title, release_date
FROM records
WHERE artist_user_id = 'user_bab3f848-261f-4056-a865-4f01793058a3'
ORDER BY release_date DESC
LIMIT 5;
```

(I tried this query on both the records and records_q5 tables. Error in both cases).

```
token@cqlsh:hoofers> SELECT record_id, title, release_date from records WHERE artist_user_id = 'user_bab3f848-261f-4056-a865-4f01793058a3' ORDER BY release_date DESC
LIMIT 5;
InvalidRequest: Error from server: code=2200 [Invalid query] message="Ordering on non-clustering column release_date requires the column to be indexed"
token@cqlsh:hoofers> SELECT record_id, title, release_date from records_q5 WHERE artist_user_id = 'user_bab3f848-261f-4056-a865-4f01793058a3' ORDER BY release_date D
ESC LIMIT 5;
InvalidRequest: Error from server: code=2200 [Invalid query] message="Ordering on non-clustering column release_date requires the column to be indexed"
token@cqlsh:hoofers> █
```

b) CQL CREATE Statement:

```
CREATE TABLE records_q6 (
    artist_user_id text,
    record_id text,
    release_date date,
    genre text,
    title text,
    PRIMARY KEY (artist_user_id, release_date, record_id)
) WITH CLUSTERING ORDER BY (release_date DESC);
```

c) Results:

```
token@cqlsh:hoofers> SELECT record_id, title, release_date from records_q6 WHERE artist_user_id = 'user_bab3f848-261f-4056-a865-4f01793058a3' LIMIT 5;
```

record_id	title	release_date
record_ff3420fe-cf7f-43f9-9131-1965883acc51	Plant worker doctor	2022-10-07
record_4f4f27a7-03f8-4adb-b96e-93adf9eb4e62	Under total throughout	2022-09-08
record_2c83cd72-4450-4936-bb9d-1732ced5a166	Money material	2022-02-18
record_8c389ed6-2101-489d-a301-b66ab43ff51c	Discover fast	2021-08-22
record_25778e66-e835-4347-9d4e-98a48f8424a1	Way real	2021-04-03

(5 rows)
token@cqlsh:hoofers> █

d) Explanation:

The new clustering key of release_date allows Cassandra to order on that column without the column being indexed. This is why the error from before is no longer raised after we include release_date as the clustering key within our CREATE TABLE statement.

7.

a) CQL Create Statement:

```
CREATE TABLE table_7a (  
    user_id text,  
    review_id text,  
    record_id text,  
    rating int,  
    PRIMARY KEY (user_id, rating, review_id) )  
WITH CLUSTERING ORDER BY (rating DESC);
```

b) CQL Create Statement:

```
CREATE TABLE table_7b (  
    genre text,  
    record_id text,  
    PRIMARY KEY (genre, record_id));
```

c) CQL Create Statement:

```
CREATE TABLE table_7c (  
    artist_user_id text,  
    posted_at timestamp,  
    review_id text,  
    record_id text,  
    title text,  
    rating int,  
    PRIMARY KEY (artist_user_id, posted_at, review_id) )  
WITH CLUSTERING ORDER BY (posted_at DESC);
```

d) CQL Create Statement:

```
CREATE TABLE table_7d (  
    user_id text,  
    initiate_at timestamp,  
    session_id text,  
    replay_count int,  
    PRIMARY KEY (user_id, initiate_at, session_id)  
);
```

8.

a)

- CQL Query:

```
SELECT review_id, record_id  
FROM table_7a WHERE user_id = 'user_9e48cbb4-0bf9-43fc-a578-213fae51068b'  
ORDER BY rating DESC LIMIT 10;
```

- Result:

```
token@cqlsh:hoofers> SELECT review_id, record_id FROM table_7a WHERE user_id = 'user_9e48cbb4-0bf9-43fc-a578-213fae51068b' ORDER BY rating DESC LIMIT 10;
```

review_id	record_id
review_23801293-5f92-4159-805f-d2a4809027f1	record_6c4b5698-4cb2-47a3-acb2-5161e51e2b47
review_23e212cf-70cf-4f5e-a151-5650e585074e	record_04b59df9-af69-423b-90c9-35bf9148835c
review_405b9a5f-22fd-433f-bf97-4fbbe7e1dd79	record_591478b0-1fa1-4f90-9094-76586f44066a
review_785f16cb-1d01-44b7-9ef1-13c4de101196	record_99f7b9ad-7969-4faf-9285-162f50c1578c
review_c0215ff1-50ca-4b07-a722-50d4d998498b	record_e4b025e5-69f2-4ef7-a2b1-8b786794c009
review_c9b4a223-5b83-4a77-897b-c35d8ffd467b	record_2c9d0def-b9e4-4577-a023-3aad8ce11806
review_fc195a51-6fff-472e-86ae-d749dad03389	record_c3f90c67-bf41-4f04-89cc-afbfa663a4cc
review_2a193df6-6a49-461e-b347-916663d656f3	record_92179aa1-4d00-4d1e-8ffb-fbbc96c196e9
review_2c28ef7e-b611-42e4-ae49-0278c5f303ac	record_c3504b5f-2504-44a7-af79-87968eaceaac
review_37df2d42-83b4-4fbc-9b30-15ee98878f8e	record_8c389ed6-2101-489d-a301-b66ab43ff51c

(10 rows)

b)

- CQL Query:

```
SELECT COUNT(*) FROM table_7b WHERE genre = 'Folk';
```

- Result:

```
token@cqlsh:hoofers> select count(*) from table_7b where genre = 'Folk';

count
-----
    57

(1 rows)
```

c)

- CQL Query:

```
SELECT review_id, record_id, title, rating
```

```
FROM table_7c WHERE artist_user_id = 'user_6f33f39e-7659-4673-bd80-ca11394424b0'
```

```
ORDER BY posted_at DESC LIMIT 10;
```

- Result:

```
token@cqlsh:hoofers> SELECT review_id, record_id, title, rating
... FROM table_7c WHERE artist_user_id = 'user_6f33f39e-7659-4673-bd80-ca11394424b0'
... ORDER BY posted_at DESC LIMIT 10;
```

review_id	record_id	title	rating
review_658cf2aa-b73f-4da6-b5f0-a9cf9edb65b5	record_1e0d5cc2-2593-4718-94a4-a87a2bc73878	Free say	3
review_346885af-b687-4d5a-82dc-e21dd6458390	record_f9220687-10b1-469c-888b-e306ed7a3376	Successful but up	2
review_3b74edf5-dc86-48ca-afe6-cd12c3677aa1	record_324b326c-9cd7-4cb7-8064-b431d3d96f4c	Police whose goal	1
review_86418842-4609-47b4-891c-c7a5c1d1ccb4	record_f274db0f-3a28-4781-9a00-dbacf5b902bf	Pick administration	4
review_85696c26-1a1b-44bc-bd61-50e5ed9a28d1	record_324b326c-9cd7-4cb7-8064-b431d3d96f4c	Police whose goal	3
review_39517c83-a010-49a4-b276-d58e423e655b	record_f0fcf759-2df7-437b-b497-b6cc4a3d726e	Hour myself seat	2
review_d828453b-13b1-420b-9921-534a8a508936	record_e793d30f-958f-499d-a8de-170732db99cd	They rather	5
review_a6a3afab-f7bc-40fd-8b71-1883e8595baa	record_f9220687-10b1-469c-888b-e306ed7a3376	Successful but up	1
review_8361944a-193c-465e-a0aa-c13349eb29e5	record_f274db0f-3a28-4781-9a00-dbacf5b902bf	Pick administration	5
review_ef32d5f6-f5cf-4816-baa5-208d3ebfd966	record_1e0d5cc2-2593-4718-94a4-a87a2bc73878	Free say	1

```
(10 rows)
token@cqlsh:hoofers>
```

d)

- CQL Query:

```
SELECT MAX(replay_count) FROM table_7d
```

```
WHERE user_id = 'user_05f9132b-47fb-4d2b-992c-17b3c4afb2df'
```

```
AND initiate_at >= '2024-08-01 00:00:00' AND initiate_at <= '2024-09-01 00:00:00';
```

- Result:

```
token@cqlsh:hoofers> SELECT MAX(replay_count) FROM table_7d
... WHERE user_id = 'user_05f9132b-47fb-4d2b-992c-17b3c4afb2df'
... AND initiate_at >= '2024-08-01 00:00:00' AND initiate_at <= '2024-09-01 00:00:00';

system.max(replay_count)
-----
4

(1 rows)
token@cqlsh:hoofers> █
```

9.

- CQL INSERT statements:

```
INSERT INTO records (record_id, artist_user_id, title, genre, release_date)
```

```
VALUES ('record_d2f498f8-d7ff-4f1c-a967-7090417751f5', 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c', 'Blue by You', 'Rock', '2024-10-07');
```

```
INSERT INTO users (user_id, city, email, genres, joined_date, nickname, state, street, zip) VALUES ('user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c', NULL, NULL, NULL, NULL, NULL, NULL, NULL, NULL);
```

```
INSERT INTO reviews (review_id, body, posted_at, rating, record_id, user_id) VALUES ('review_new', NULL, NULL, NULL, 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5', 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c');
```

```
INSERT INTO sessions (session_id, device, initiate_at, leave_at, music_quality, record_id, remaining_time, replay_count, track_number, user_id) VALUES ('session_new', NULL, NULL, NULL, NULL, 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5', NULL, NULL, NULL, 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c');
```

```
INSERT INTO table_7a (user_id, rating, review_id, record_id) VALUES ('user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c', 5, 'review_new', 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5');
```

```
INSERT INTO table_7b (genre, record_id) VALUES ('Rock', 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5');
```

```
INSERT INTO table_7c (artist_user_id, posted_at, review_id, rating, record_id, title) VALUES ('user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c', '2024-10-08', 'review_new', 5, 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5', 'Blue By You');
```

```
INSERT INTO table_7d (user_id, initiate_at, session_id, replay_count) VALUES ('user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c', '2024-10-08T11:00:00Z', 'session_new', 1);
```

- Verification queries:

SELECT * FROM records WHERE record_id = 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5';

SELECT * FROM users WHERE user_id = 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c';

SELECT * FROM reviews WHERE review_id = 'review_new';

SELECT * FROM sessions WHERE session_id = 'session_new';

SELECT * FROM table_7a WHERE user_id = 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c';

SELECT * FROM table_7b WHERE record_id = 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5';

SELECT * FROM table_7c WHERE record_id = 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5';

SELECT * FROM table_7d WHERE user_id = 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c';

- Result:

First 4 queries above.

```
token@cqlsh:hoofers> SELECT * FROM records WHERE record_id = 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5';
```

record_id	artist_user_id	genre	release_date	title
record_d2f498f8-d7ff-4f1c-a967-7090417751f5	user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c	Rock	2024-10-07	Blue by You

(1 rows)

```
token@cqlsh:hoofers> SELECT * FROM users WHERE user_id = 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c';
```

user_id	city	email	genres	joined_date	nickname	state	street	zip
user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c	null	null	null	null	null	null	null	null

(1 rows)

```
token@cqlsh:hoofers> SELECT * FROM reviews WHERE review_id = 'review_new';
```

review_id	body	posted_at	rating	record_id	user_id
review_new	null	null	null	record_d2f498f8-d7ff-4f1c-a967-7090417751f5	user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c

(1 rows)

```
token@cqlsh:hoofers> SELECT * FROM sessions WHERE session_id = 'session_new';
```

session_id	device	initiate_at	leave_at	music_quality	record_id	remaining_time	replay_count	track_number	user_id
session_new	null	null	null	null	record_d2f498f8-d7ff-4f1c-a967-7090417751f5	null	null	null	user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c

(1 rows)

```
token@cqlsh:hoofers>
```

Verifying table_7a

```
token@cqlsh:hoofers> SELECT * FROM table_7a WHERE user_id = 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c' AND rating = 5 AND review_id='review_new';
```

user_id	rating	review_id	record_id
user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c	5	review_new	record_d2f498f8-d7ff-4f1c-a967-7090417751f5

(1 rows)

```
token@cqlsh:hoofers>
```

Verifying table_7b

```
token@cqlsh:hoofers> SELECT * FROM table_7b WHERE record_id = 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5' ALLOW FILTERING;
```

genre	record_id
Rock	record_d2f498f8-d7ff-4f1c-a967-7090417751f5

(1 rows)

```
token@cqlsh:hoofers>
```

Verifying table_7c

```
token@cqlsh:hoofers> SELECT * FROM table_7c WHERE record_id = 'record_d2f498f8-d7ff-4f1c-a967-7090417751f5' ALLOW FILTERING;
```

artist_user_id	posted_at	review_id	rating	record_id	title
user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c	2024-10-08 00:00:00.000000+0000	review_new	5	record_d2f498f8-d7ff-4f1c-a967-7090417751f5	Blue By You

```
(1 rows)
token@cqlsh:hoofers>
```

Verifying table_7d

```
token@cqlsh:hoofers> SELECT * FROM table_7d WHERE user_id = 'user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c' AND initiate_at = '2024-10-08T11:00:00Z' AND session_id = 'session_new';
```

user_id	initiate_at	session_id	replay_count
user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c	2024-10-08 11:00:00.000000+0000	session_new	1

```
(1 rows)
token@cqlsh:hoofers>
```

10.

Python script:

```
from cassandra.cluster import Cluster
```

```
from cassandra.query import SimpleStatement
```

```
from cassandra.auth import PlainTextAuthProvider
```

```
from datetime import datetime
```

```
# Connect to your Cassandra cluster
```

```
# Path to your secure connect bundle
```

```
secure_connect_bundle = '/Users/ameya/Desktop/DataLoading/secure-connect-cs224p-fall.zip'
```

```
# Astra DB credentials
```

```
# For client_id and client_secret, I have not filled them out here, but in my code I would fill out the  
# necessary values from the downloaded JSON token file.
```

```
client_id = 'client_id'
```

```
client_secret = 'client_secret'
```

```
# Set up authentication and connection
```

```
auth_provider = PlainTextAuthProvider(client_id,client_secret)
```

```
cluster = Cluster(cloud={'secure_connect_bundle': secure_connect_bundle},  
auth_provider=auth_provider)
```

```
session = cluster.connect('hoofers')
```

This function replicates the query from question 9 to add the record to all the tables.

It takes as input all the possible parameters that may be of interest, but we might only use a few of them per table

The '%s' serves as a place holder for the relevant variables for each table.

```
def addRecord(session, record_id, artist_user_id, title, genre, release_date, user_id, review_id,
rating, session_id):
```

```
    # Insert into records table
```

```
    session.execute( SimpleStatement( "INSERT INTO records (record_id, artist_user_id,
title,      genre,      release_date)  VALUES      (%s,      %s,      %s,      %s,      %s)",
consistency_level=ConsistencyLevel.ONE ), (record_id, artist_user_id, title, genre, release_date)
)
```

```
    # Insert into users table
```

```
    session.execute( SimpleStatement( "INSERT INTO users (user_id, city, email, genres,
joined_date, nickname, state, street, zip) VALUES (%s, NULL, NULL, NULL, NULL, NULL, NULL,
NULL, NULL)", consistency_level=ConsistencyLevel.ONE ), (user_id,) )
```

```
    # Insert into reviews table
```

```
    session.execute( SimpleStatement( "INSERT INTO reviews (review_id, body, posted_at,
rating,  record_id,  user_id)  VALUES      (%s,  NULL,  NULL,  %s,  %s,  %s)",
consistency_level=ConsistencyLevel.ONE ), (review_id, rating, record_id, user_id) )
```

```
    # Insert into sessions table
```

```
    session.execute( SimpleStatement( "INSERT INTO sessions (session_id, device,
initiate_at, leave_at, music_quality, record_id, remaining_time, replay_count, track_number,
user_id) VALUES (%s, NULL, NULL, NULL, NULL, %s, NULL, NULL, NULL, %s)",
consistency_level=ConsistencyLevel.ONE ), (session_id, record_id, user_id) )
```

```
    # Insert into table_7a
```

```
    session.execute( SimpleStatement( "INSERT INTO table_7a (user_id, rating, review_id,
record_id) VALUES (%s, %s, %s, %s)", consistency_level=ConsistencyLevel.ONE ), (user_id,
rating, review_id, record_id) )
```

```
    # Insert into table_7b
```

```
    session.execute( SimpleStatement( "INSERT INTO table_7b (genre, record_id) VALUES
```

```
(%s, %s)", consistency_level=ConsistencyLevel.ONE ), (genre, record_id) )
```

```
# Insert into table_7c
```

```
session.execute( SimpleStatement( "INSERT INTO table_7c (artist_user_id, posted_at,
review_id, rating, record_id, title) VALUES (%s, %s, %s, %s, %s, %s)",
consistency_level=ConsistencyLevel.ONE ), (artist_user_id, datetime.now().isoformat(),
review_id, rating, record_id, title) )
```

```
# Insert into table_7d
```

```
session.execute( SimpleStatement( "INSERT INTO table_7d (user_id, initiate_at,
session_id, replay_count) VALUES (%s, %s, %s, %s)",
consistency_level=ConsistencyLevel.ONE ), (user_id, datetime.now().isoformat(), session_id, 1)
)
```

```
# Query to call the function and add the given record to the database, connected to 'hoofers'
#keyspace.
```

```
session = connect_to_cassandra()
```

```
add_record( session,
```

```
'record_632fe768-eeeb-4596-9780-cc21734feec5',
```

```
'user_b91cf915-487b-42fc-b6b8-6c17935bb755',
```

```
'One Sour Day',
```

```
'R&B',
```

```
'2024-10-07',
```

```
'user_b91cf915-487b-42fc-b6b8-6c17935bb755',
```

```
'review_new',
```

```
5,
```

```
'session_new' )
```

```
session.shutdown()
```