

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import torch
import torch.nn as nn
import torch.optim as optim

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

from torch.utils.data import DataLoader, TensorDataset

In [ ]: from sklearn.datasets import load_iris

# Load the Iris dataset
iris = load_iris()

# You can now access the data and target variables
x = iris.data
y = iris.target

# You can also see the feature names and target names
feature_names = iris.feature_names
target_names = iris.target_names

print("Features (X):")
print(X[15]) # Print the first 5 rows of features

print("\nTargets (y):")
print(y[15]) # Print the first 5 target values

print("\nFeature Names:")
print(feature_names)
print("\nTarget Names:")
print(target_names)

Features (X):
[[5.1 3.5 1.4 0.2]
 [4.9 3. 1.4 0.2]
 [4.7 3.2 1.3 0.2]
 [4.6 3.1 1.5 0.2]
 [5. 3.6 1.4 0.2]]

Targets (y):
[0 0 0 0 0]

Feature Names:
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']

Target Names:
['setosa' 'versicolor' 'virginica']
```

Data visualization and pre-processing

The PCA plot helps illustrate overall separability between classes.

```
In [ ]: # Create a DataFrame from the data
df = pd.DataFrame(X, columns=feature_names)
df['species'] = y # Add target column

# Optional: Replace numerical labels with actual species names
df['species'] = df['species'].map({i: name for i, name in enumerate(target_names)})

In [ ]: print(df.head())
print(df.describe())
print(df['species'].value_counts())

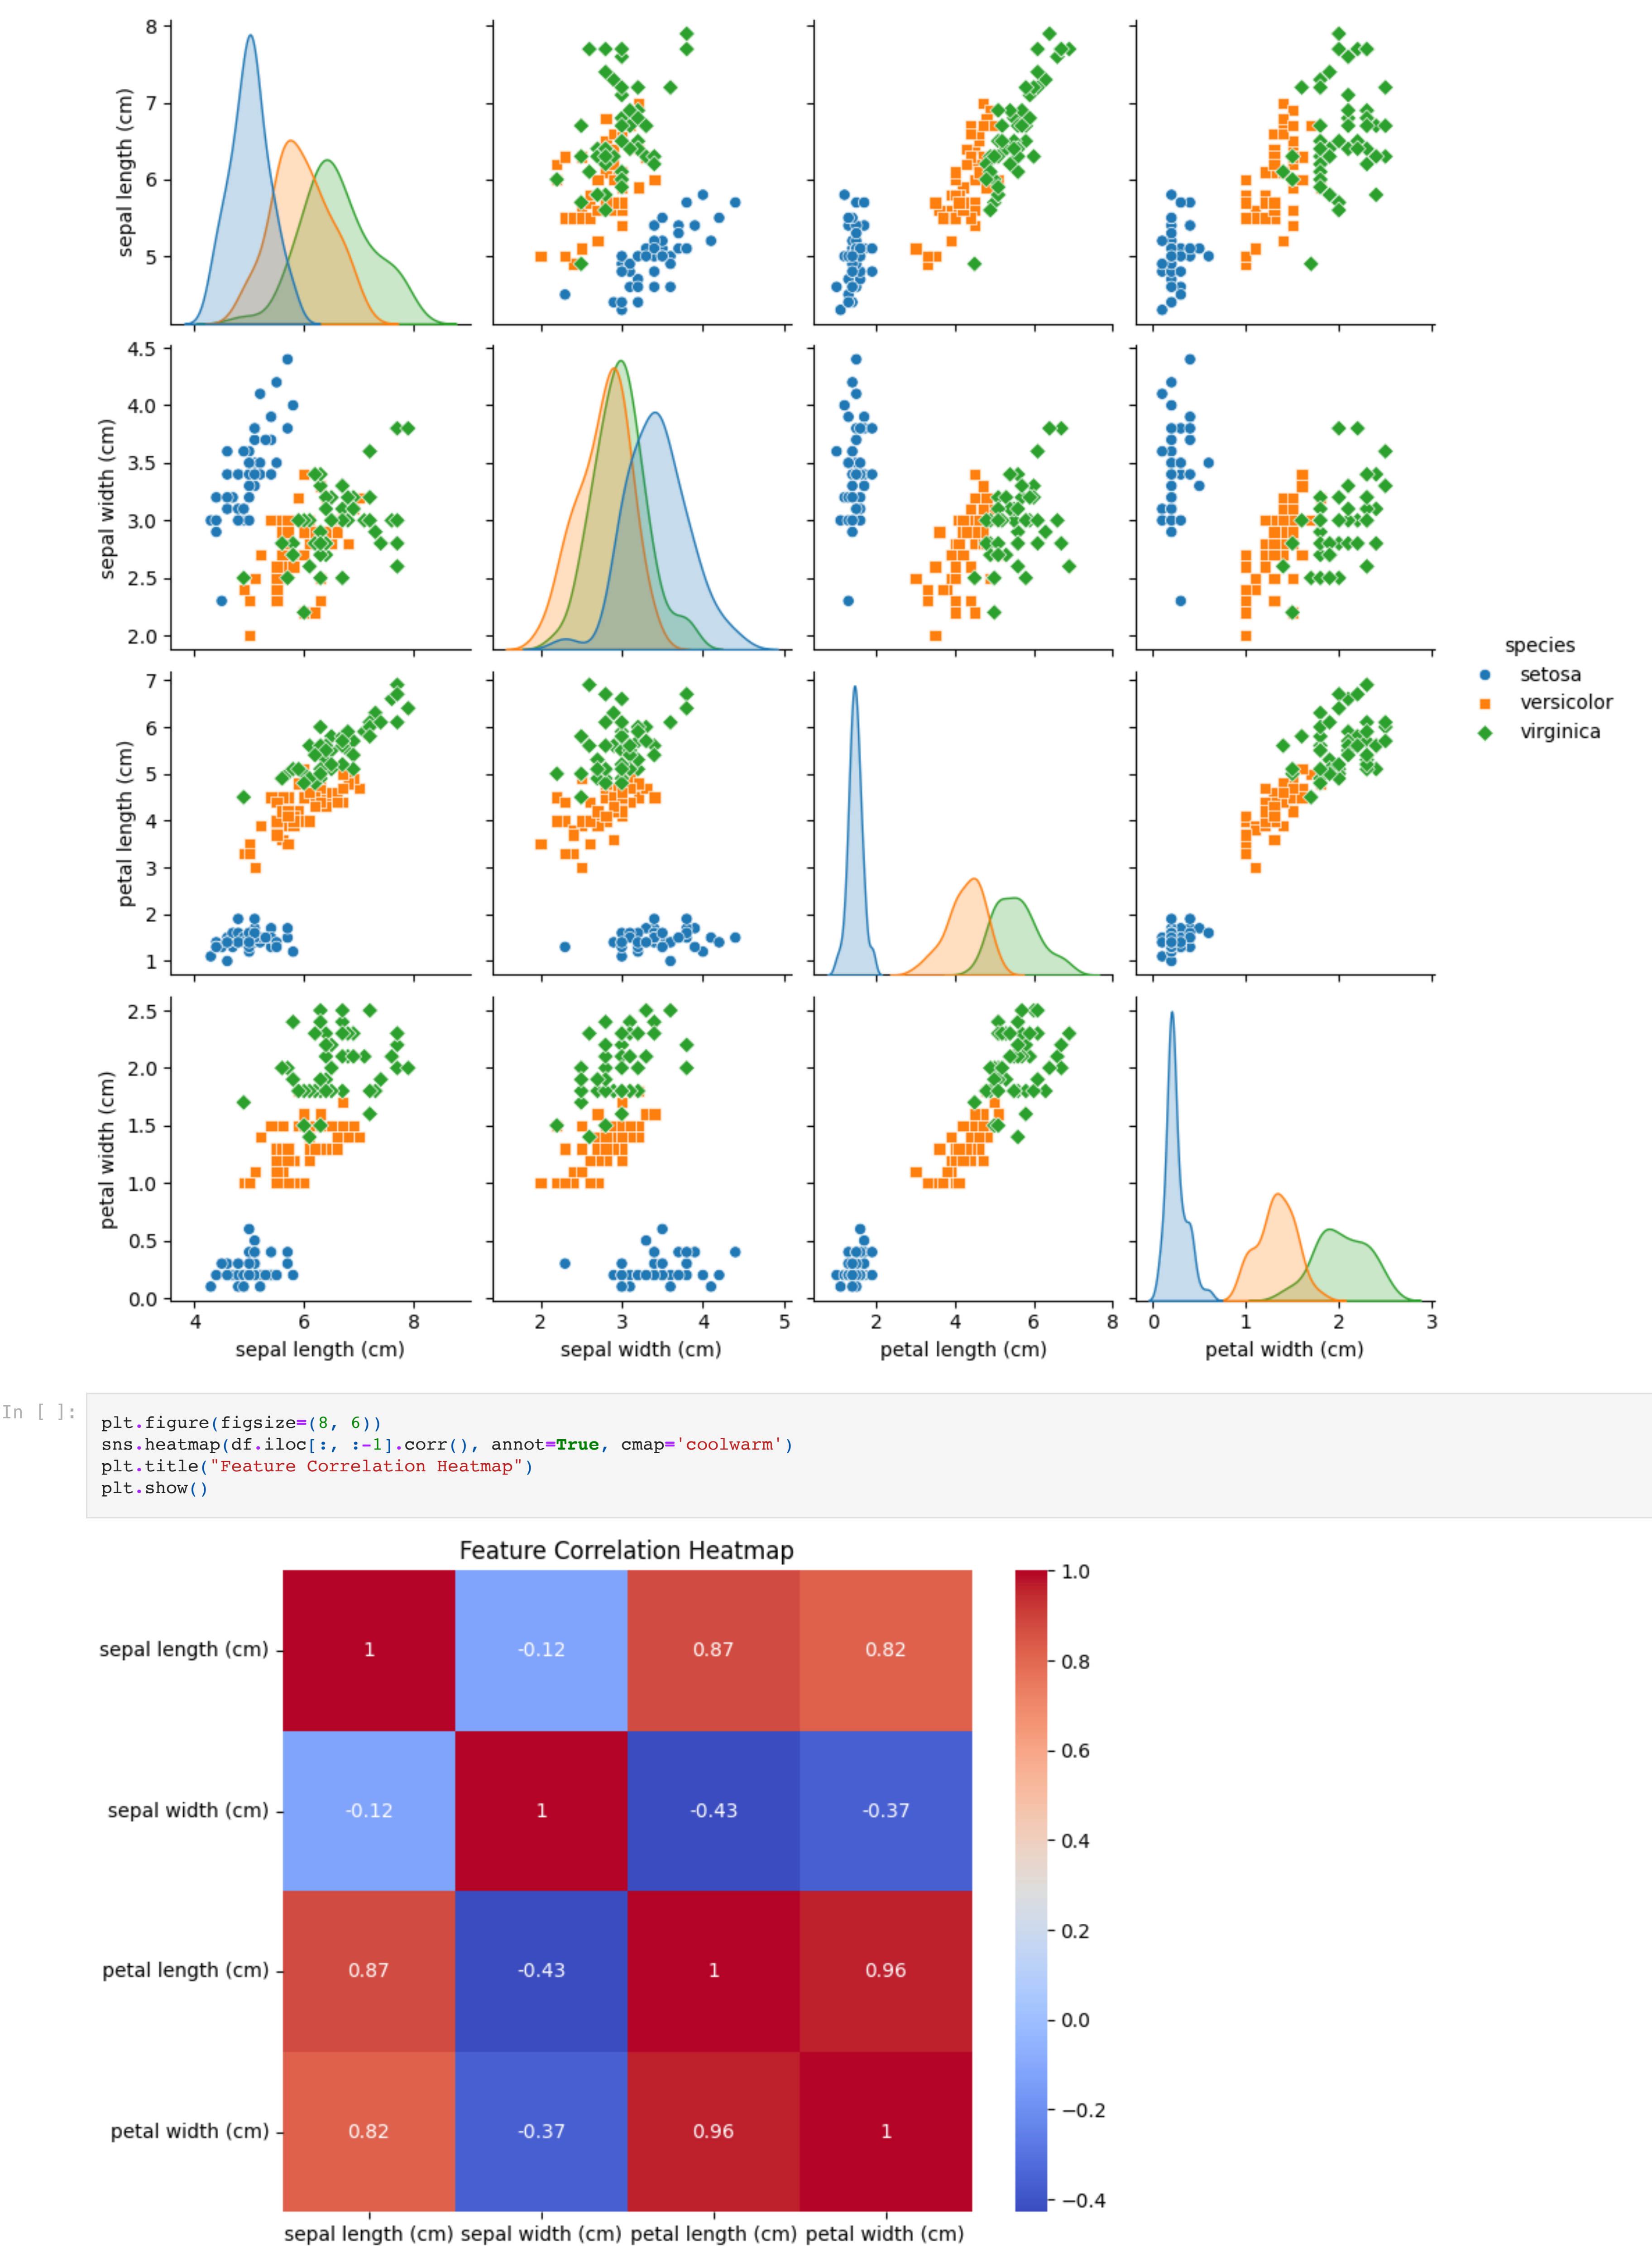
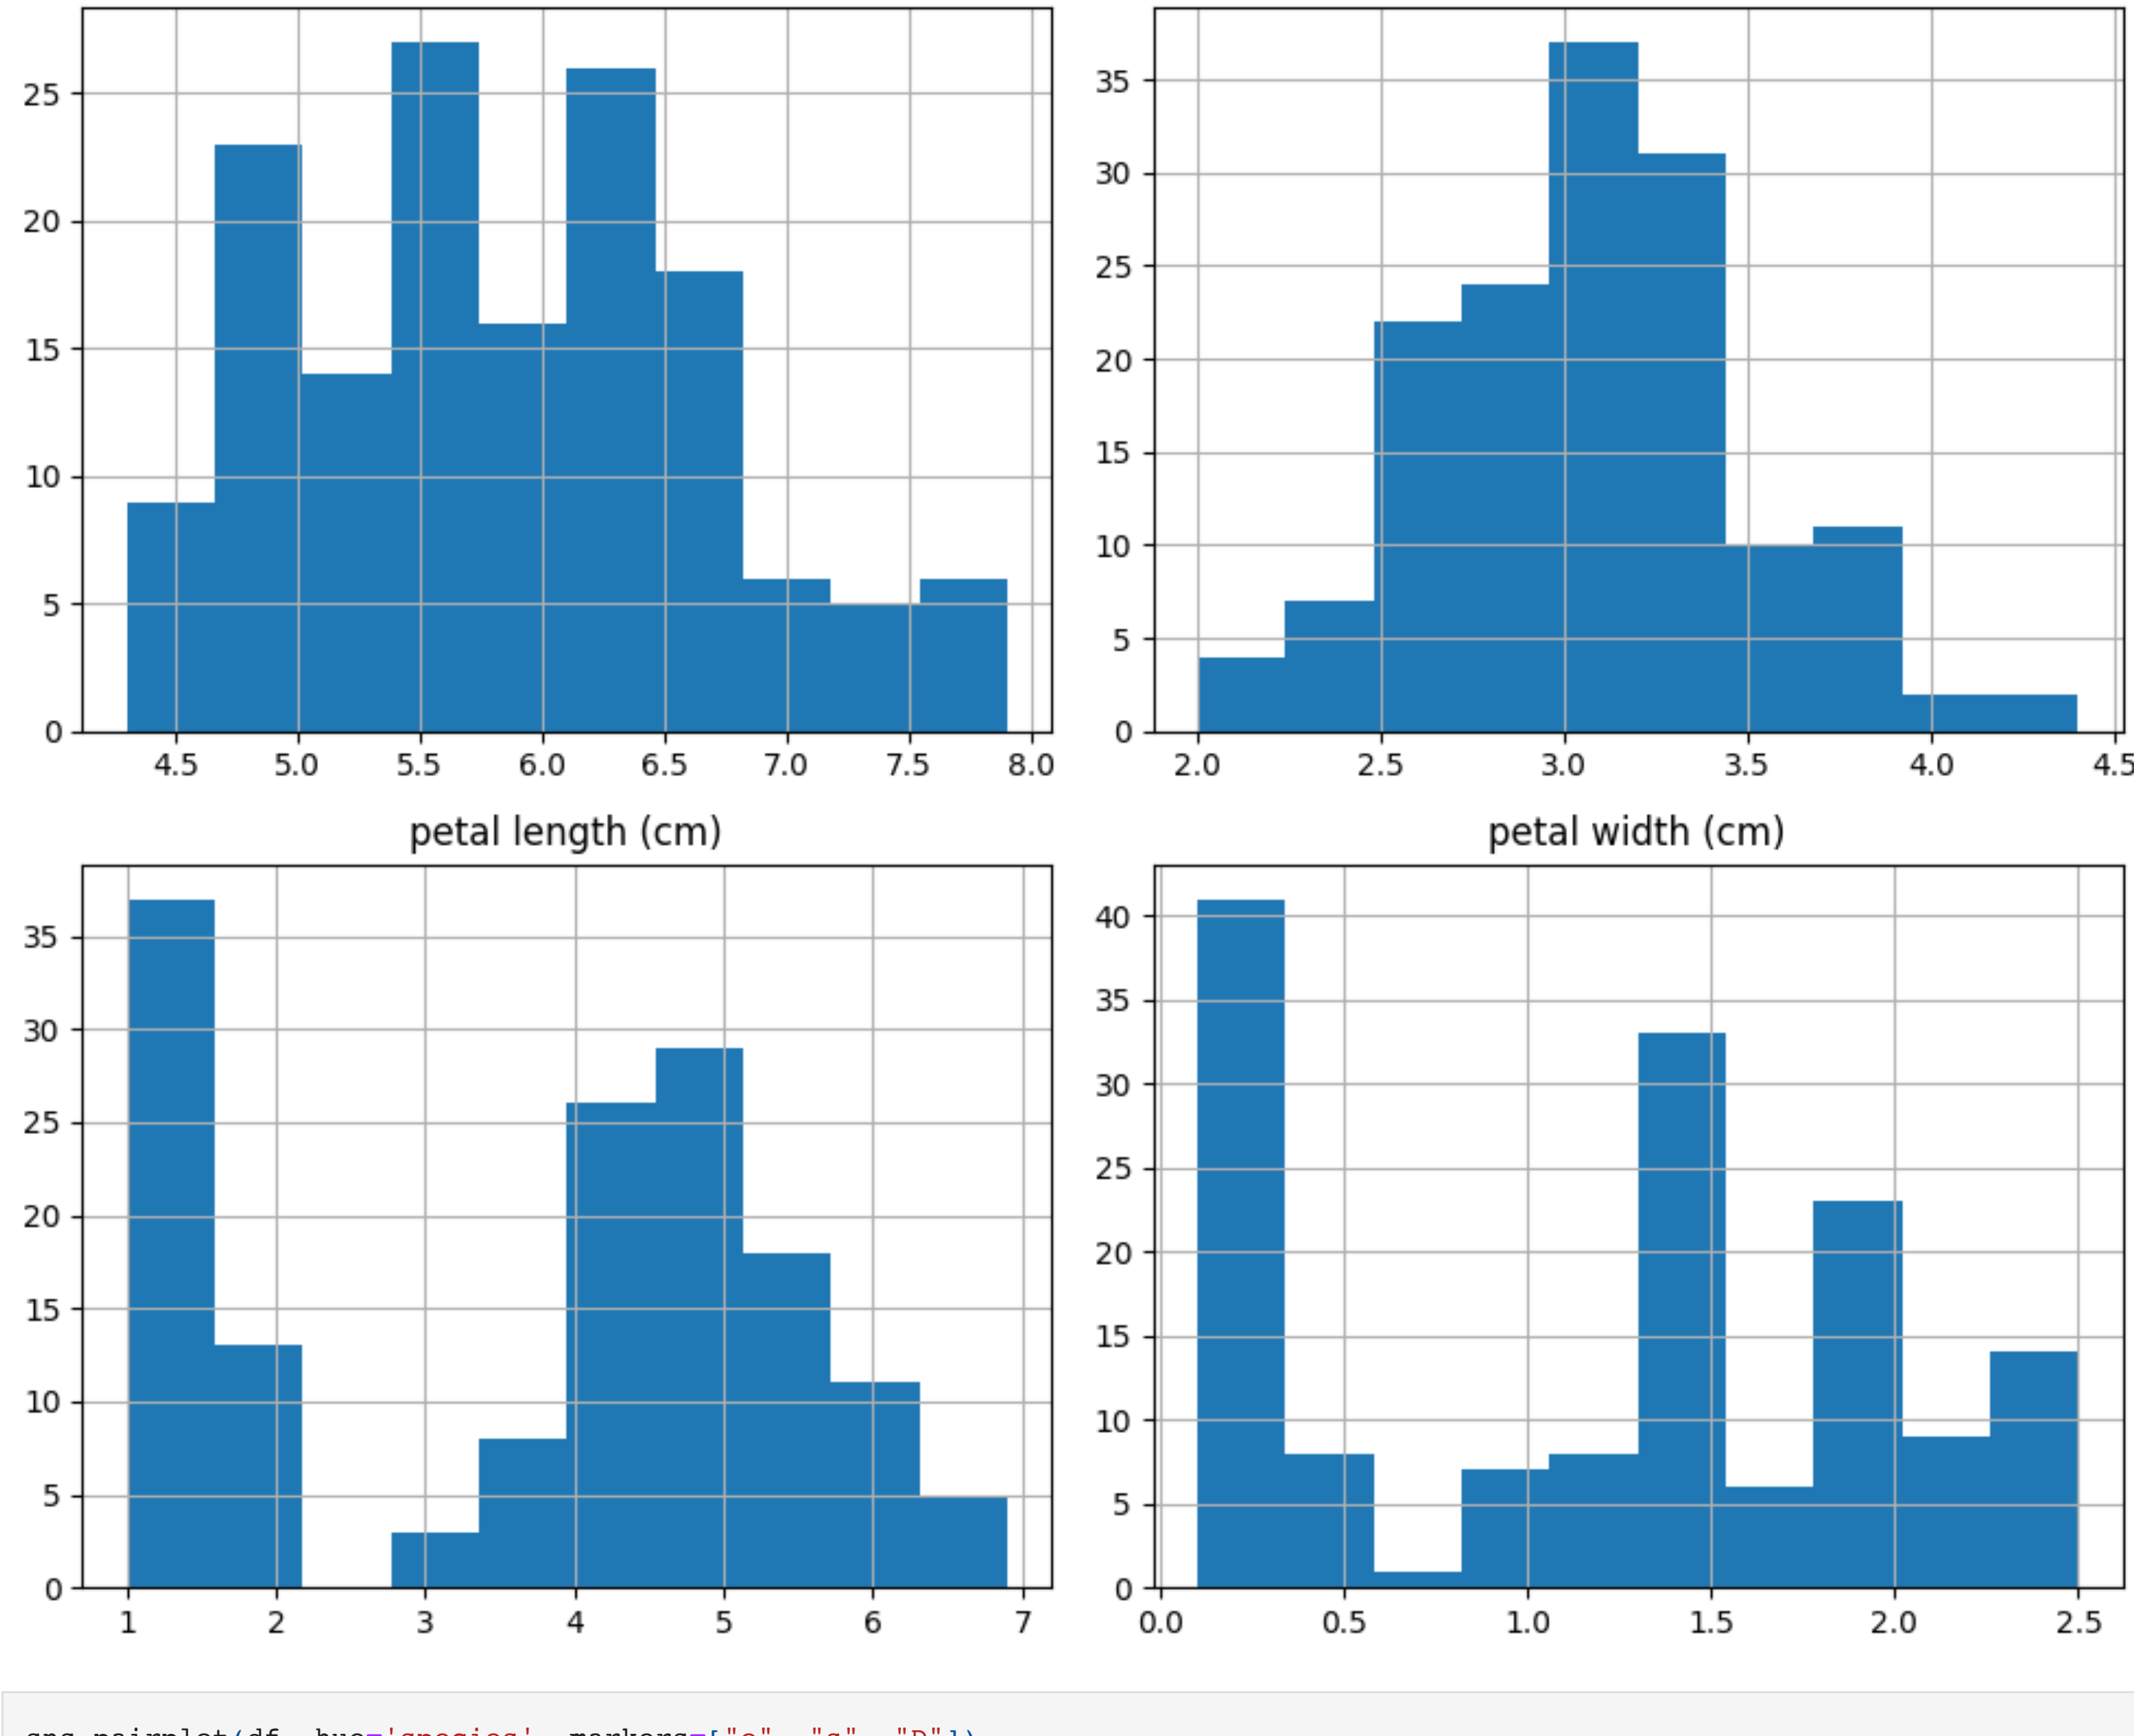
   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm) \
0              5.1              3.5              1.4              0.2
1              4.9              3.0              1.4              0.2
2              4.7              3.2              1.3              0.2
3              4.6              3.1              1.5              0.2
4              5.0              3.6              1.4              0.2

   species
0  setosa
1  setosa
2  setosa
3  setosa
4  setosa

   sepal length (cm)  sepal width (cm)  petal length (cm) \
count      150.000000      150.000000      150.000000
mean         5.843333         3.057333         3.758000
std          0.828066         0.435866         1.765298
min          4.300000         2.000000         1.000000
25%          5.100000         2.800000         1.600000
50%          5.800000         3.000000         4.350000
75%          6.400000         3.300000         5.100000
max          7.900000         4.400000         6.900000

   petal width (cm)
count      150.000000
mean         1.199333
std          0.762238
min          0.100000
25%          0.300000
50%          1.300000
75%          1.800000
max          2.500000
species
setosa      50
versicolor  50
virginica   50
Name: count, dtype: int64

In [ ]: df.hist(figsize=(10, 8))
plt.tight_layout()
plt.show()
```



Feature Correlation Heatmap

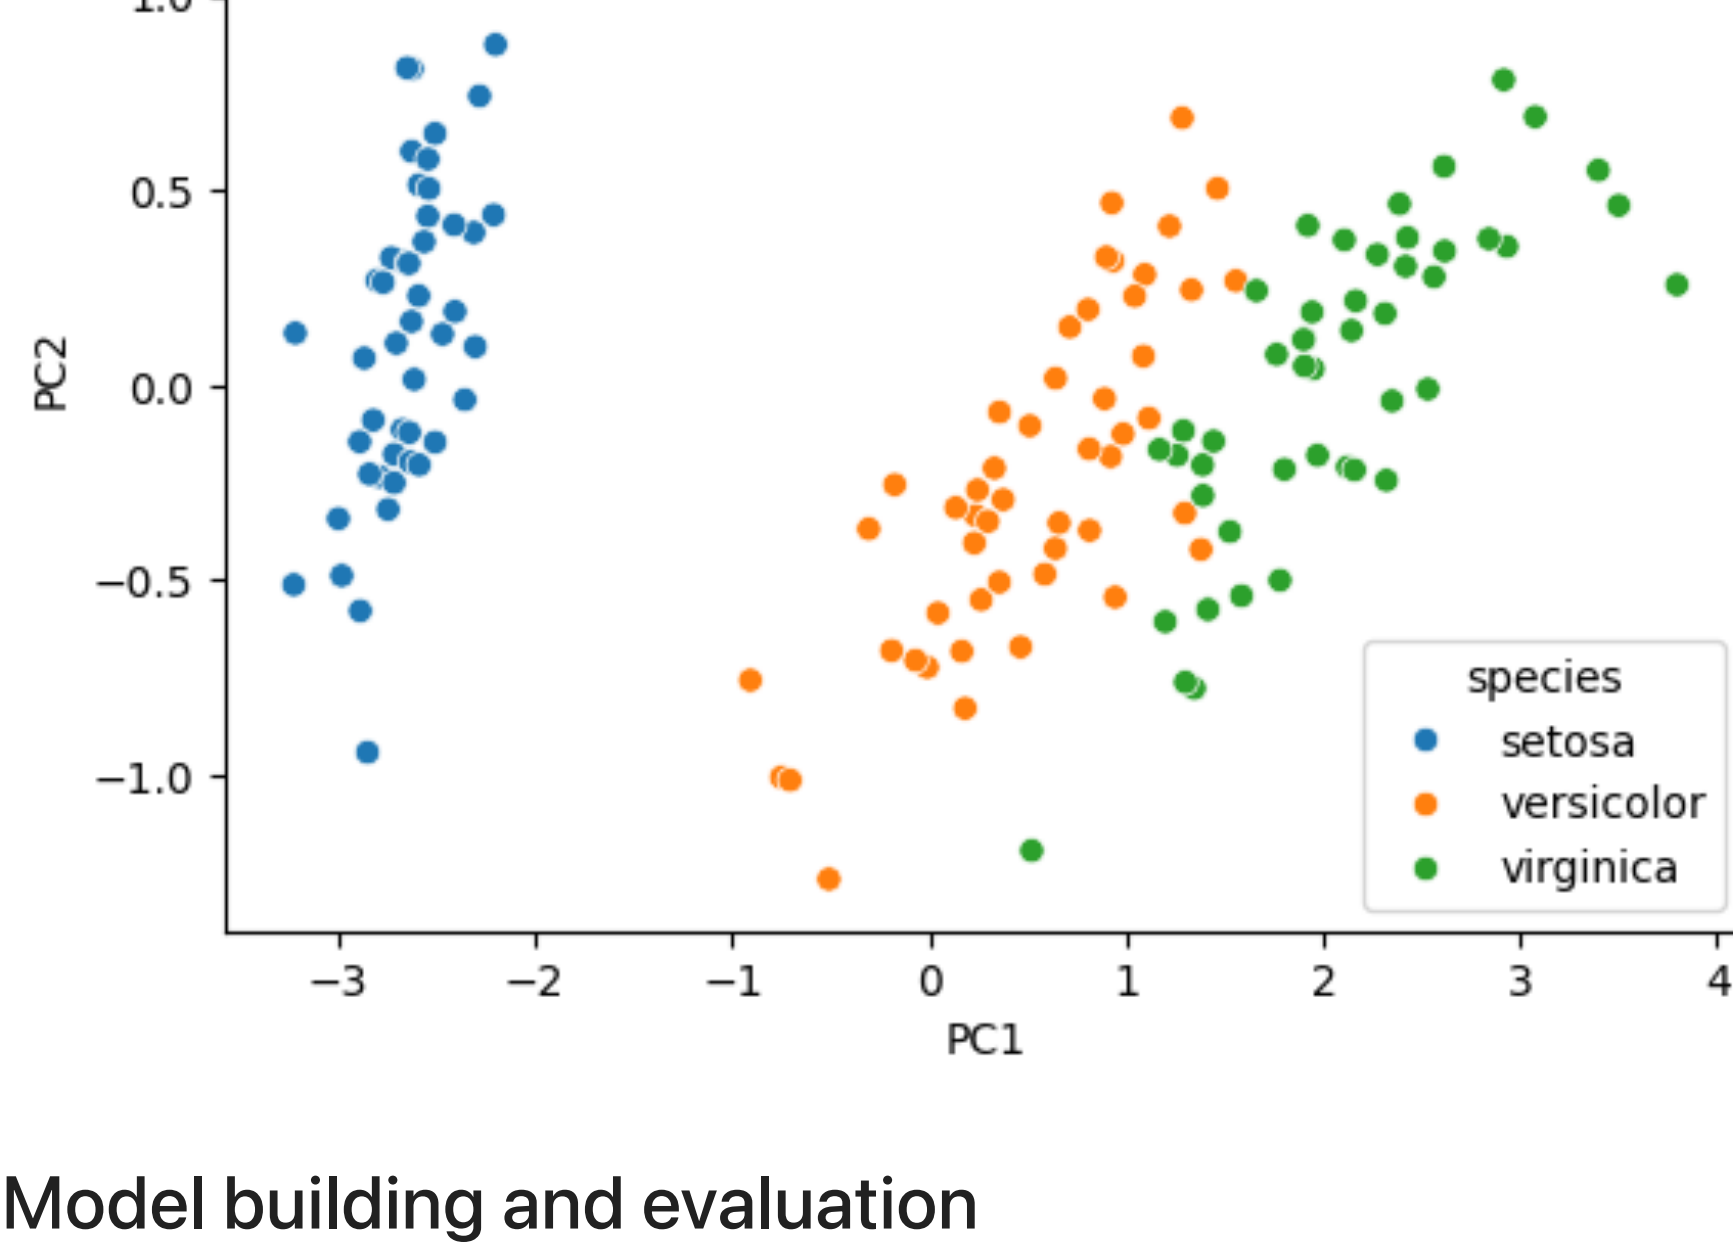


```
In [ ]: from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

df_pca = pd.DataFrame(X_pca, columns=['PC1', 'PC2'])
df_pca['species'] = df['species']

sns.scatterplot(data=df_pca, x='PC1', y='PC2', hue='species')
plt.title("PCA of Iris Dataset")
plt.show()
```



## Model building and evaluation

StandardScaler() standardizes the data, which means each feature is transformed to have mean = 0 and variance = 1. Even when original features do not follow Normal distribution, standardization is helpful because features will have zero mean and be on the same scale with unit variance. So, gradient descent will converge faster.

Can consider MinMaxScaler if data is too skewed or has many outliers.

Also, it is a good idea to shuffle training data to avoid overfitting and speed up convergence. Without shuffling, model might adapt to patterns in the training data and poorly generalize.

```
In [ ]: # Standardize features (mean = 0, variance = 1).
# Makes them zero-centered with same scale.
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-test split (80% train, 20% test)

# Note: STRATIFICATION is necessary only if output classes are imbalanced. Ex,
# if we have 90% class A and 10% class B, we want these proportionally represented
# in the train and test sets.

x_train, x_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Convert to PyTorch tensors

# PyTorch models work with tensors. Specify float32 datatype for input, and long
# datatype for output.
x_train_tensor = torch.tensor(x_train, dtype=torch.float32)
x_test_tensor = torch.tensor(x_test, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.long)
y_test_tensor = torch.tensor(y_test, dtype=torch.long)

# Create datasets and dataloaders

# Dataset wrapper allows pairing inputs and targets.
# PyTorch knows how to fetch a sample as: (X_train[i], y_train[i])
train_dataset = TensorDataset(x_train_tensor, y_train_tensor)
test_dataset = TensorDataset(x_test_tensor, y_test_tensor)

# DataLoader splits dataset into mini-batches (here, 5 samples per batch)
# Optionally shuffles the data (important during training!)

train_loader = DataLoader(train_dataset, batch_size=5, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=5)

# Don't shuffle test_loader, so we can consistently evaluate model.
```

## Model architecture

We define the architecture with 2 hidden layers, each with 10 neurons. All layers are fully-connected (fc). The final layer is the output layer with 3 neurons, each representing an output class. ReLU is our activation function to introduce some non-linearity into the model.

```
In [ ]: # nn.Module is base class for all neural network modules in PyTorch
class IrisNet(nn.Module):
    def __init__(self):
        super(IrisNet, self).__init__()

        # nn.Linear applies a linear transformation (y = xA + bias) to input
        self.fc1 = nn.Linear(4, 10) # Input layer - Hidden layer
        self.fc2 = nn.Linear(10, 10) # Hidden layer - Hidden layer
        self.fc3 = nn.Linear(10, 3) # Hidden layer - Output layer (3 classes)

    def forward(self, x):
        x = torch.relu(self.fc1(x)) # Activation after first layer
        x = torch.relu(self.fc2(x)) # Activation after second layer
        x = self.fc3(x) # Output logits (no softmax here)
        return x

# Instantiate the model
model = IrisNet()
print(model)

IrisNet(
  (fc1): Linear(in_features=4, out_features=10, bias=True)
  (fc2): Linear(in_features=10, out_features=10, bias=True)
  (fc3): Linear(in_features=10, out_features=3, bias=True)
)
```

## Loss and optimizer

Our loss function is CrossEntropyLoss. You don't have to apply softmax yourself in your model because CrossEntropyLoss automatically applies log(softmax(...)) to the model's output.

So no need to add softmax layer to output.

Suppose model output (logits) are [1.2, 0.3, -0.8], and the true class label is 0 (from 0, 1, 2).

Softmax then converts to probabilities [0.65, 0.25, 0.10], so we are saying P(y=0) = 0.65. The cross-entropy loss is -log(0.65) = 0.43.. If the model was more confident (e.g., predicted 0.95 for class 0), the loss would be lower.

So the cross-entropy loss is determined by how confident the model is that the sample falls in true class.

```
In [ ]: # Loss function: CrossEntropyLoss is ideal for multi-class classification
criterion = nn.CrossEntropyLoss()

# Optimizer: Adam is adaptive and works well with default settings
optimizer = optim.Adam(model.parameters(), lr=0.01)
```

## Training loop

```
In [ ]: num_epochs = 50

for epoch in range(num_epochs):
    model.train() # Set model to training mode
    running_loss = 0.0

    for batch_X, batch_y in train_loader:

        optimizer.zero_grad() # Clear gradients
        outputs = model(batch_X) # Forward pass

        # Gives average loss per sample in that batch.
        loss = criterion(outputs, batch_y) # Compute loss
        loss.backward() # Backpropagation
        optimizer.step() # Update weights

        running_loss += loss.item() # Accumulate running loss for each batch

    if (epoch + 1) % 10 == 0:
        avg_loss = running_loss / len(train_loader)
        print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}')

Epoch [10/50], Loss: 0.0768
Epoch [20/50], Loss: 0.0545
Epoch [30/50], Loss: 0.0547
Epoch [40/50], Loss: 0.0535
Epoch [50/50], Loss: 0.0687
```

## Evaluation on test set

```
In [ ]: model.eval() # Set model to evaluation mode
correct = 0
total = 0

with torch.no_grad(): # Disable gradient computation for efficiency
    for batch_X, batch_y in test_loader:
        outputs = model(batch_X) # Forward pass

        _, predicted = torch.max(outputs, 1) # Get predicted class index
        total += batch_y.size(0) # Update total sample count
        correct += (predicted == batch_y).sum().item() # Count correct predictions

accuracy = 100 * correct / total
print(f'Test Accuracy: {accuracy:.2f}%')

Test Accuracy: 96.67%
```

## Possible next steps

- Try dropout or batch normalization

```
In [ ]:
```